




Revealing Potential Biases in LLM-Based Recommender Systems in the Cold Start Setting

Alexandre Andre *, Gauthier Roy *, Eva Dyer , Kai Wang 



University of Pennsylvania



Georgia Institute of Technology

* Equal Contribution



RecSys 2025
Prague


The 2nd Workshop on Evaluating and Applying Recommender Systems with Large Language Models (EARL) at RecSys 2025.

**What university
should I go to?**

AI LLM

- 1. University A**
- 2. University B**
- 3. University C**

Cold Start Setting Implications





Cold start

New user:

- No/low context.
- Still access to some information.
- Maybe **sensitive** information (the logging information).
- **Critical moment** $\Rightarrow$ fair access + discoverability


Automated Pipeline


- Plug & play pipeline of evaluation $\Rightarrow$ benchmark




How to Evaluate the B ias of a LLM-Based Recommender Systems in the Cold Start Setting?


How to Evaluate *Bias*?





“What  would you recommend to a  User” (Neutral)




How to Evaluate *Bias*?




“What  would you recommend to a  User” (Att. A)



Neutral




How to Evaluate *Bias*?



$$\textit{Bias}(\text{Att. A}) = 1 - \text{sim}\left(\begin{array}{c} \text{Harry Potter} \\ \text{Avengers} \\ \text{T-Rex} \end{array}, \begin{array}{c} \text{Harry Potter} \\ \text{Avengers} \\ \text{Rings} \end{array} \right)$$

Different Attributes $\Rightarrow$ Different *Bias*?




Examples with Sim = IOU

$$Bias(\text{Att. A}) = 1 - IOU(\text{Att. A}, \text{Neutral}) = 1 - \frac{|\text{Intersection}|}{|\text{Union}|} = 1 - \frac{|\begin{array}{|c|} \hline \text{Glasses} \quad \text{No Glasses} \\ \hline \end{array}|}{|\begin{array}{|c|} \hline \text{Glasses} \quad \text{No Glasses} \quad \text{T-Rex} \quad \text{No T-Rex} \\ \hline \end{array}|} = 1 - \frac{2}{4} = 0.5$$



$$Bias(\text{Att. B}) = 1 - IOU(\text{Att. B}, \text{Neutral}) = 1 - \frac{|\text{Intersection}|}{|\text{Union}|} = 1 - \frac{|\begin{array}{|c|} \hline \text{Glasses} \\ \hline \end{array}|}{|\begin{array}{|c|} \hline \text{Glasses} \quad \text{No Glasses} \quad \text{T-Rex} \quad \text{No Glasses} \quad \text{No T-Rex} \\ \hline \end{array}|} = 1 - \frac{1}{5} = 0.8$$

Also other Sim: SERP and PRAG




How to Flexibly Automate the Evaluation ?

How to Flexibly Automate the Evaluation ?



What is happening here?

Some Prompt Engineering



Input prompt

Context prompt

Consider the following list of **{dataset type}**: **{items}**. You will recommend exactly **{k}** items chosen ONLY from this exact list.

+

Query prompt

Based on the list provided above, recommend an ORDERED list of exactly **{k}** **{dataset type}** for **{sensitive attribute}** **{type of activity}**. I cannot provide more preferences due to privacy.

+

Example prompt

--- EXAMPLE --- If the provided list started with: **{example input}**... And you were asked for **{k}** items, your output should look EXACTLY like this (using actual recommendations) **{example output}** --- END EXAMPLE ---

+


Format instruction prompt

IMPORTANT: FINAL OUTPUT FORMATTING: 1. Your response MUST contain ONLY the final ordered list of **{k}** recommended **{dataset type}**. 2. Format the list as NUMBERS followed by a PERIOD and a SPACE, then the item name, with each item on a NEW LINE. Like the example above. 3. CRITICAL: Use the EXACT item names as provided in the input list. DO NOT abbreviate, add descriptions (like '(MIT)' or ' - *...*'), change capitalization, or paraphrase. 4. ABSOLUTELY NO introductory text (like 'Okay, here is the list...'), concluding remarks, explanations, apologies, reasoning, notes, or any text whatsoever before or after the numbered list. ONLY THE LIST.

System prompt


You are an expert recommender system for **{dataset type}**. Your ONLY task is to generate an ordered list of recommendations based on the user query. You MUST follow all formatting instructions precisely.

+



How to Use the Pipeline?

Datasets




Movie Database (TMDB) API





Spotify Web API



2023 QS World University Rankings (Kaggle)





Sensitive Attributes




Main categories:

- Gender
- Country
- Parents occupation
- Continent
- Race
- Religion



Models




Llama 3.2 3B: meta-llama/Llama-3.2-3B-Instruct 🤗

Gemma 3 4B: google/gemma-3-4b-it 🤗

Dataset	College🎓		Music🎵		Movie🎬	
Model	Gemma	Llama	Gemma	Llama	Gemma	Llama
IOU	.55 ± .12	.50 ± .14	.27 ± .09	.46 ± .15	.55 ± .09	.54 ± .11
SERP	.83 ± .03	.80 ± .05	.77 ± .02	.76 ± .05	.81 ± .03	.80 ± .04
PRAG	.47 ± .09	.46 ± .12	.20 ± .06	.37 ± .14	.46 ± .08	.50 ± .13

Gemma 3 {1/4/12}B: google/gemma-3-{1/4/12}b-it 🤗

Hypotheses to Investigate



- **H1:** **Larger** LLMs exhibit less bias

$$Bias_{\text{Large model}}(X) < Bias_{\text{Model}}(X)$$

- **H2:** LLMs replicate **societal** stereotypes

$$Bias_{\text{Model}}(X) \approx Bias_{\text{Society}}(X)$$


- **H3:** Adding **context** to a user mitigates bias

$$Bias_{\text{Model}}(X|\text{Context}) < Bias_{\text{Model}}(X)$$

- **H4:** LLMs are biased towards **Western** content


$$Bias_{\text{Model}}(X_{\text{Western}}) < Bias_{\text{Model}}(X_{\overline{\text{Western}}})$$

H1: Larger LLMs exhibit less bias



- H1: **Larger** LLMs exhibit less bias

$Bias_{\text{Large model}}(X)$



$Bias_{\text{Model}}(X)$


Dataset	College 			Music 			Movie 		
Model Size	1B	4B	12B	1B	4B	12B	1B	4B	12B
IOU	.59 ± .12	.55 ± .12	.65 ± .07	.73 ± .20	.27 ± .09	.66 ± .08	.78 ± .08	.55 ± .09	.75 ± .07
SERP	.83 ± .03	.84 ± .03	.86 ± .03	.84 ± .14	.77 ± .02	.85 ± .03	.86 ± .04	.81 ± .03	.89 ± .03
PRAG	.46 ± .10	.48 ± .09	.55 ± .09	.69 ± .23	.20 ± .06	.48 ± .09	.72 ± .09	.46 ± .08	.67 ± .07

Non-linear relationship between scale and bias.

Explanation:

- Big model $\Rightarrow$ follows “too well” instructions (over-emphasize sensitive attributes).
- Small model $\Rightarrow$ does not follow instruction (bad re-ranking + not staying within the dataset item list).
- Good balance for 4B.

H2: LLMs replicate societal stereotypes



- H2: LLMs replicate **societal** stereotypes

$$Bias_{\text{Model}}(x)$$




$$Bias_{\text{Society}}(x)$$


Societal biases observed.

Explanation:

- Top-20 movie recommendations generated by Gemma 3 4B.
- *Neutral* result overlaps with *boy/male* not with *girl/female*.
- Even more divergence for Gemma 3 12B.



H3: Adding context to a user mitigates bias




- **H3:** Adding **context** to a user mitigates bias

$$Bias_{Model}(X|Context) \leq Bias_{Model}(X)$$


Context reduced bias.

Explanation:

- Movie recommendations with Gemma 3 12B.
- Context $\Rightarrow$ less bias.
- Context $\Rightarrow$ more uniform bias within category.
- Aligned with platform effort to collect user preference in cold start.



H4: LLMs are biased towards Western content




- **H4:** LLMs are biased towards **Western** content

$$Bias_{Model}(X_{Western}) \leq Bias_{Model}(X_{\text{---}})$$


Western bias observed.

Explanation:

- Movie recommendations with Gemma 3 12B.
- Neutral $\Rightarrow$ a lot of Western content
- Neutral aligned with , ,  ...
- Bias in the setting (English language) + American model (not Qwen/DeepSeek).




Hypotheses Results



- **H1:** **Larger** LLMs exhibit less bias

$$Bias_{\text{Large model}}(X)$$



$$Bias_{\text{Model}}(X)$$

- **H2:** LLMs replicate **societal** stereotypes

$$Bias_{\text{Model}}(X)$$



$$Bias_{\text{Society}}(X)$$

- **H3:** Adding **context** to a user mitigates bias

$$Bias_{\text{Model}}(X|Context)$$



$$Bias_{\text{Model}}(X)$$


- **H4:** LLMs are biased towards **Western** content

$$Bias_{\text{Model}}(X_{\text{Western}})$$



$$Bias_{\text{Model}}(X_{\text{Western}})$$

Conclusions



Main contributions

Benchmark & Pipeline: **modular**, **extensible** benchmark for cold-start LLM recommendations supporting any:

- Open-source LLM
- Dataset
- Attribute

Discussion

Complex relationship between **bias** and **personalization** :

When does it become unfair stereotyping and risk reinforcing societal biases?

